


Lecture 6

File Systems and Pattern Matching

Troy Pretty
Digital Forensic Analyst

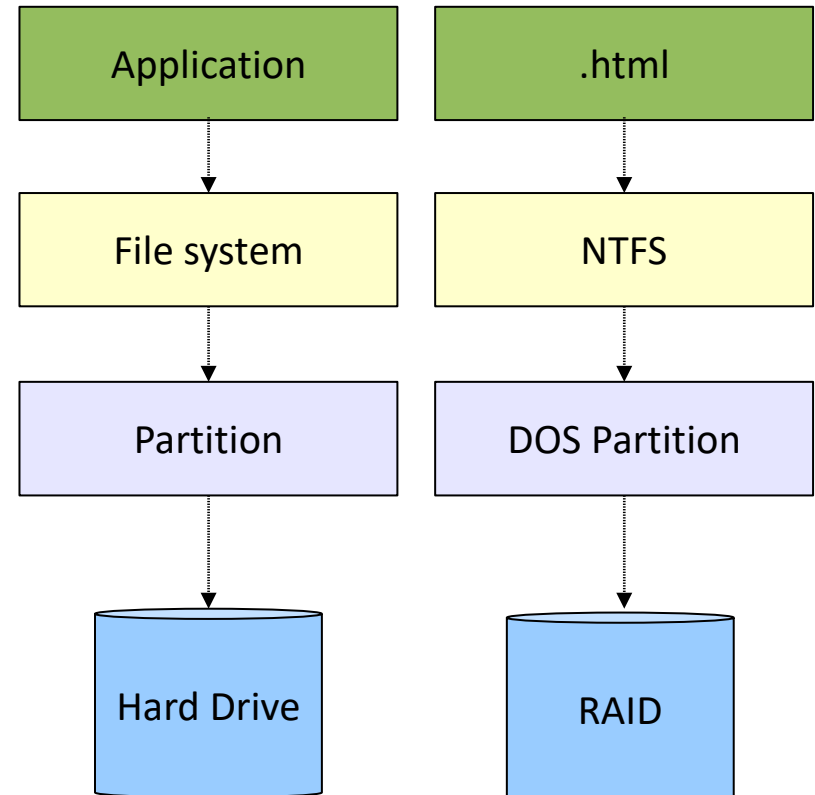


Outline and Learning Goals

- General features of file systems
- FAT and NTFS
- Searching file systems
- File deletion vs. file erasing

File Systems

- A file system is a way of organising and store data on a disk according to well defined standards
- Imposes a structure on the partition scheme



Widely used file systems

- FAT

- File Allocation Table
- Simple and one of the earliest widely used file systems
- Still widely used in flash memory devices such as thumb drives and digital cameras
- Design goal is simplicity
- Limits:
 - FAT12 – 16MB Volume
 - FAT16 – 4GB Volume
 - FAT32 – 2TB Volume, 4GB File Size

Widely used file systems

- exFAT
 - Extended File Allocation Table
 - Introduced around the time of Windows Vista
 - Designed to be lightweight like FAT but without the extra overhead of NTFS
 - Overcomes volume and file size limitations of FAT

Widely used file systems

- NTFS

- New Technologies File System
- Default system for Microsoft Windows since NT
- One of the more complex file systems
- Design goal is flexibility
 - Security
 - Compression
 - Encryption

Widely used file systems

- EXT

- Extended File System (EXT2, EXT3, EXT4)
- Default file system for most Linux distributions
- Based on UFS (but considerably simplified)
- Design goals are speed and reliability

Widely used file systems

- HFS/HFS+

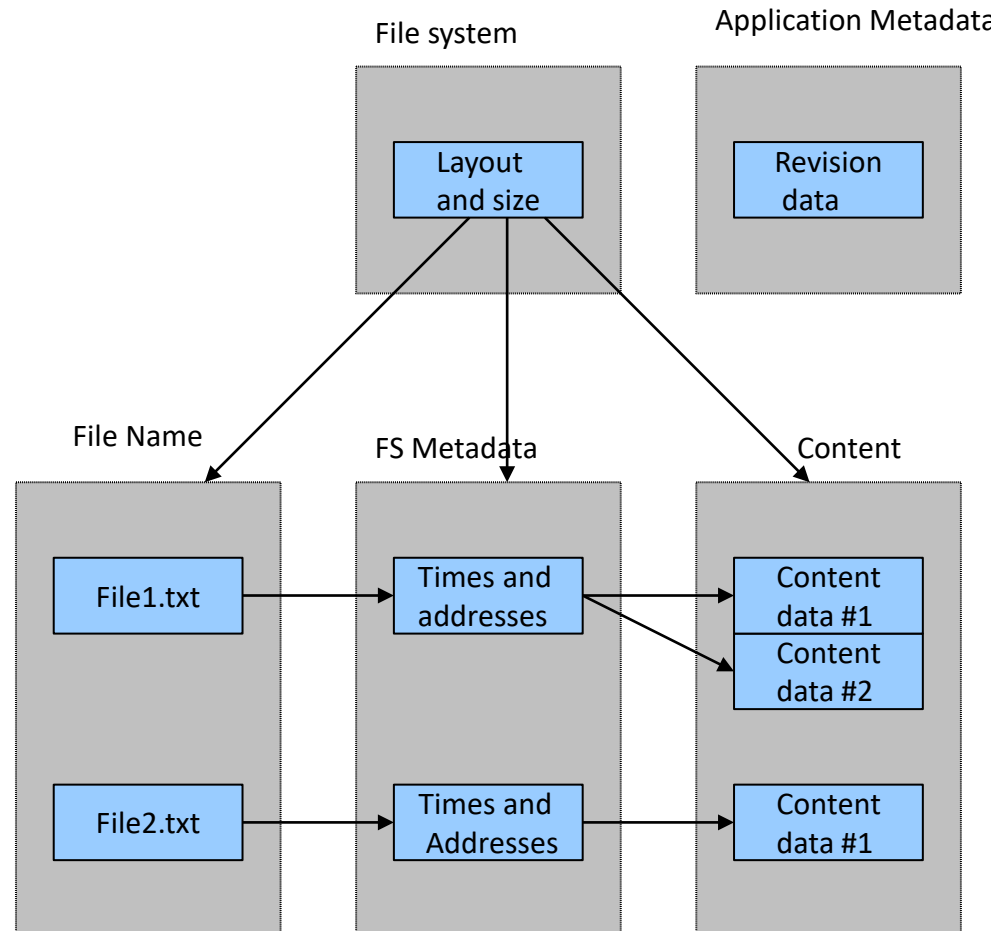
- Hierarchical File System/Hierarchical File System Plus
- Apple proprietary filesystem
- Used in OSX and on Apple iDevices
- Linus Torvalds the creator of the Linux kernel has been quoted to say:
 - “Quite frankly, HFS+ is probably the worst file-system ever”

- APFS

- Apple File System
- Replaces HFS+
- Universal across all Apple devices
- First released on iPhone running iOS 10.3 (March 2017)
- Increase in speed and better disk utilisation

File systems

- The major file systems have some common features
 - File system information – A map of the file system
 - File name – Human readable label
 - FS Metadata – Data that describes files
 - Content – Data stored in file system
 - Application Metadata – Specific information depending on file type



File systems

- Analysis of a file system is usefully carried out according to the categories described in the previous slide
 - File system information, content, FS metadata, file name, application metadata
- Depending on what we are searching for, one particular category might be more appropriate as the basis of the search than the others

File system category

- A map of the file system
- Enables identification of the location of other important data
- Usually located in a standard data structure in the first few sectors of the system
- An analysis of a file system will usually start here

File system category

- File system is FAT32
- Volume name is PENDRIVE
- Boot sector is 0
- Data area is sector 7840 to 4004319
- Obtained using the fsstat command

```
philip@philip-G41M-Combo: ~/Desktop
File Edit View Search Terminal Help
philip@philip-G41M-Combo:~/Desktop$ fsstat -f fat fat.dd | more
FILE SYSTEM INFORMATION
-----
File System Type: FAT32

OEM Name: MSDOS5.0
Volume ID: 0x6864dea4
Volume Label (Boot Sector): NO NAME
Volume Label (Root Directory): PENDRIVE
File System Type Label: FAT32
Next Free Sector (FS Info): 26000
Free Sector Count (FS Info): 2201232

Sectors before file system: 32

File System Layout (in sectors)
Total Range: 0 - 4004319
* Reserved: 0 - 33
** Boot Sector: 0
** FS Info Sector: 1
** Backup Boot Sector: 6
* FAT 0: 34 - 3936
* FAT 1: 3937 - 7839
* Data Area: 7840 - 4004319
** Cluster Area: 7840 - 4004319
*** Root Directory: 7840 - 7847

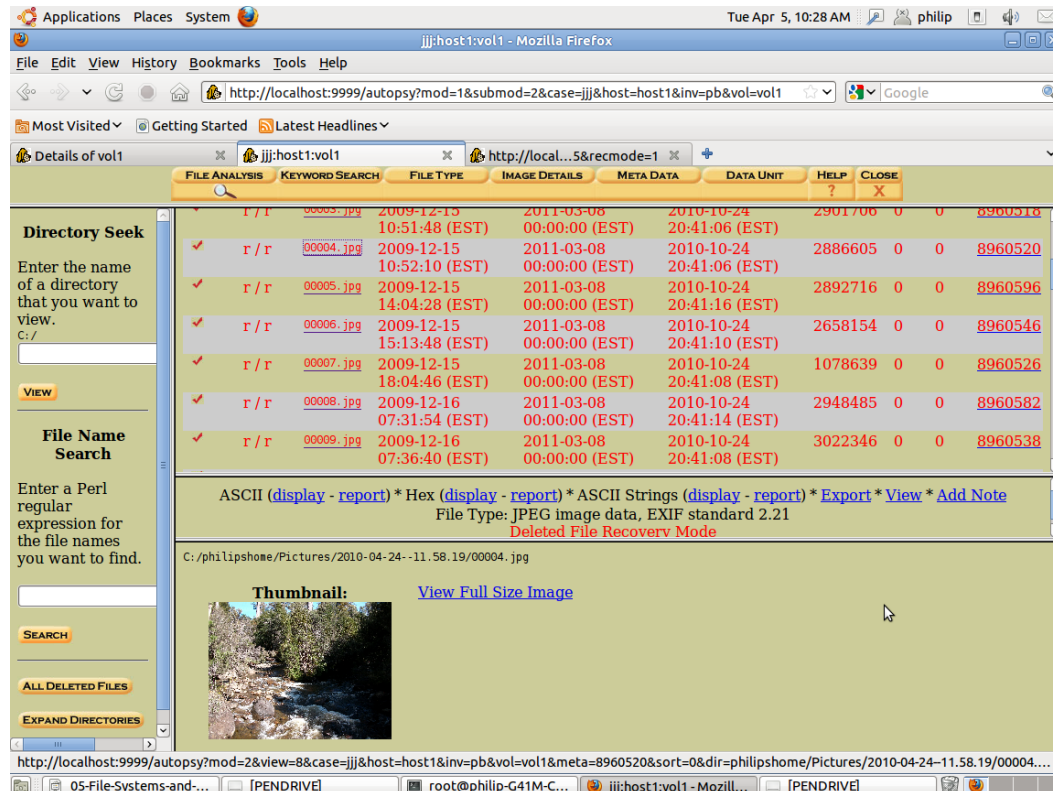
METADATA INFORMATION
-----
Range: 2 - 63943686
Root Directory: 2

CONTENT INFORMATION
-----
Sector Size: 512
Cluster Size: 4096
Total Cluster Range: 2 - 499561

FAT CONTENTS (in sectors)
-----
7840-7847 (8) -> EOF
8384-8831 (448) -> EOF
8832-9415 (584) -> EOF
9416-15431 (6016) -> EOF
```

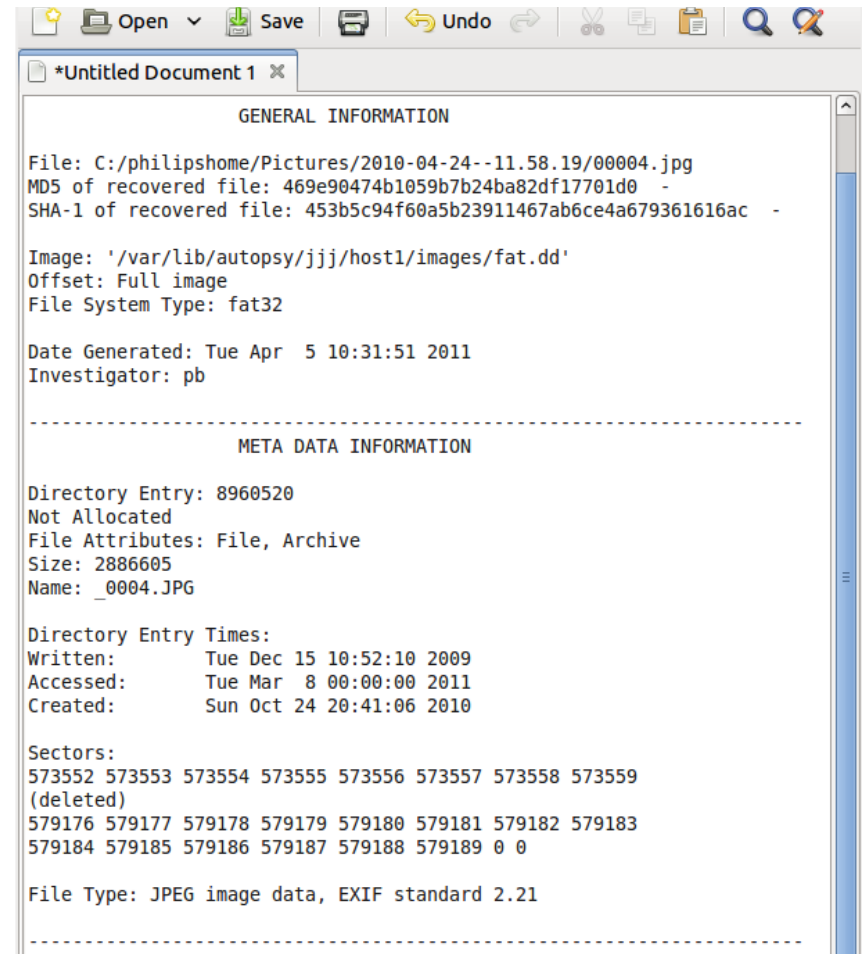
Content category

- File and directory contents
- Figure shows a deleted file from fat.dd obtained using Sleuth Kit / Autopsy



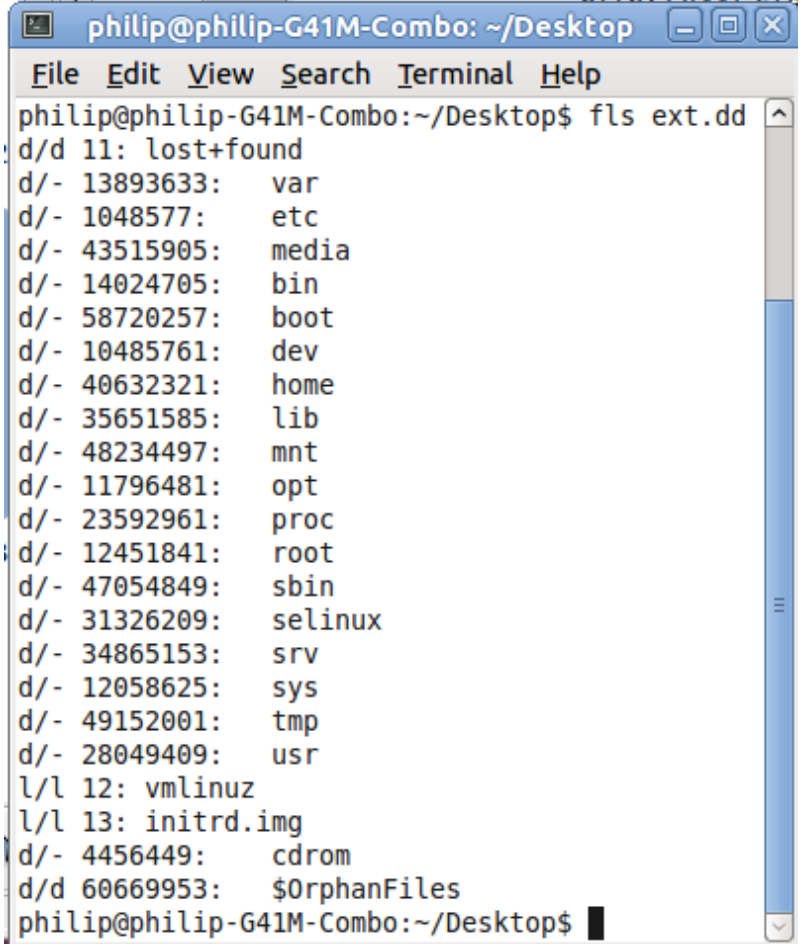
Metadata category

- Metadata from deleted file obtained using TSK/Autopsy
- Shows attributes of file _0004.jpg



Filename Category

- Associates a name with a filesystem entry
- In example we see directory /var associated with address 13803633
 - An ext file system running Linux



```
philip@philip-G41M-Combo: ~/Desktop
File Edit View Search Terminal Help
philip@philip-G41M-Combo:~/Desktop$ ls -l
d/d 11: lost+found
d/- 13893633: var
d/- 1048577: etc
d/- 43515905: media
d/- 14024705: bin
d/- 58720257: boot
d/- 10485761: dev
d/- 40632321: home
d/- 35651585: lib
d/- 48234497: mnt
d/- 11796481: opt
d/- 23592961: proc
d/- 12451841: root
d/- 47054849: sbin
d/- 31326209: selinux
d/- 34865153: srv
d/- 12058625: sys
d/- 49152001: tmp
d/- 28049409: usr
l/l 12: vmlinuz
l/l 13: initrd.img
d/- 4456449: cdrom
d/d 60669953: $OrphanFiles
philip@philip-G41M-Combo:~/Desktop$
```

Application category


- Interprets file at the application level

File Analysis Window

File Type: JPEG image data, EXIF standard 2.21

Deleted File Recovery Mode

File Name	File Type	File Size	File Date	File Time	File Location
00003.jpg	JPEG image data, EXIF standard 2.21	2901706	2009-12-15	10:51:48 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00004.jpg	JPEG image data, EXIF standard 2.21	2886605	2009-12-15	10:52:10 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00005.jpg	JPEG image data, EXIF standard 2.21	2892716	2009-12-15	14:04:28 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00006.jpg	JPEG image data, EXIF standard 2.21	2658154	2009-12-15	15:13:48 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00007.jpg	JPEG image data, EXIF standard 2.21	1078639	2009-12-15	18:04:46 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00008.jpg	JPEG image data, EXIF standard 2.21	2948485	2009-12-16	07:31:54 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg
00009.jpg	JPEG image data, EXIF standard 2.21	3022346	2009-12-16	07:36:40 (EST)	C:/philipshome/Pictures/2010-04-24-11.58.19/00004.jpg

Thumbnail:  View Full Size Image

FAT concepts and analysis

- File Allocation Table (FAT) one of the simplest file systems in common use
 - Used in older versions of Windows but main use now is with flash memory systems such as thumb drives and digital camera cards
- Each file and directory is allocated a directory entry that contains the file name, size, starting address of the file (and additional metadata)
- File and directory content is stored in disk clusters

FAT system layout

- Reserved area
 - Boot sector
- FAT area
 - Location of File Allocation Tables
 - Links to root directory entries
 - May be two FAT tables for reliability

```
root@philip-G41M-Combo: /home/philip/Desktop
File Edit View Search Terminal Help
root@philip-G41M-Combo:/home/philip/Desktop# fsstat
FILE SYSTEM INFORMATION
-----
File System Type: FAT32

OEM Name: MSDOS5.0
Volume ID: 0x6864dea4
Volume Label (Boot Sector): NO NAME
Volume Label (Root Directory): PENDRIVE
File System Type Label: FAT32
Next Free Sector (FS Info): 26000
Free Sector Count (FS Info): 2201232

Sectors before file system: 32

File System Layout (in sectors)
Total Range: 0 - 4004319
* Reserved: 0 - 33
** Boot Sector: 0
** FS Info Sector: 1
** Backup Boot Sector: 6
* FAT 0: 34 - 3936
* FAT 1: 3937 - 7839
* Data Area: 7840 - 4004319
** Cluster Area: 7840 - 4004319
*** Root Directory: 7840 - 7847

METADATA INFORMATION
-----
Range: 2 - 63943686
Root Directory: 2

CONTENT INFORMATION
-----
Sector Size: 512
Cluster Size: 4096
Total Cluster Range: 2 - 499561

FAT CONTENTS (in sectors)
-----
7840-7847 (8) -> EOF
8384-8831 (448) -> EOF
```

FAT system layout

- FAT table consists of entries for each file
- Size of each entry specified by FAT type
 - FAT12 has 12 bits, FAT16 has 16, FAT32 32 bits
 - In FAT32 28 bits are used to address the starting cluster
- Root directory
 - Link to clusters in data area
 - In FAT12/16 preallocated
 - In FAT32 part of data structure (so no bounds on size)
- Data area

FAT system root directory

- Root directory contains
 - File name (up to 11 characters) in FFFFFFFF.XXX format
 - Long file names stored in separate entries
 - File attributes such as RW, hidden, label, whether it has a long file name, whether it is a directory
 - Time / date created/written/accessed
 - First cluster address of the file

Recovering deleted files in FAT systems

- It is often possible to recover deleted files in a FAT system
- When a file is deleted the system only makes a deletion mark on the file's directory entry
- Clusters in the DATA area are marked as empty in the directory BUT are not immediately overwritten
 - When writing new data the original file clusters might be overwritten
 - But often the old data is still available

FAT deleted file recovery in TSK/Autopsy


The screenshot displays the Autopsy web interface for file recovery. The top navigation bar includes links for 'Most Visited', 'Getting Started', and 'Latest Headlines'. The main interface is divided into several sections:

- Directory Seek:** A sidebar on the left with a text input field and a 'VIEW' button.
- File Name Search:** A sidebar on the left with a text input field and a 'SEARCH' button.
- Current Directory:** Displays the path 'C:/ /philipshome/ /Pictures/ /2010-04-24--11.58.19/'. Below this are buttons for 'ADD NOTE' and 'GENERATE MD5 LIST OF FILES'.
- Table:** A table listing files with columns: DEL, Type, NAME, WRITTEN, ACCESSED, CREATED, SIZE, UID, GID, and META. The table contains five entries, including directories and files like '00001.jpg', '00002.jpg', and '00003.jpg'.
- File Details:** A section below the table showing details for the selected file '00004.jpg', including its ASCII and Hex representations, file type ('JPEG image data, EXIF standard 2.21'), and a 'Deleted File Recovery Mode' indicator.
- Thumbnail:** A section at the bottom showing a thumbnail of the selected file and a link to 'View Full Size Image'.

DEL	Type	NAME	WRITTEN	ACCESSED	CREATED	SIZE	UID	GID	META
✓	d / d	../	2010-04-26 12:31:32 (EST)	2010-10-24 00:00:00 (EST)	2010-10-24 20:41:16 (EST)	4096	0	0	942495
✓	d / d	./	2011-03-20 04:32:40 (EST)	2011-03-20 00:00:00 (EST)	2010-10-24 20:41:16 (EST)	4096	0	0	5855250
✓	r / r	00001.jpg	2009-12-04 13:29:56 (EST)	2011-03-08 00:00:00 (EST)	2010-10-24 20:41:08 (EST)	3162665	0	0	8960534
✓	r / r	00002.jpg	2009-12-04 13:36:00 (EST)	2011-03-08 00:00:00 (EST)	2010-10-24 20:41:14 (EST)	3063448	0	0	8960586
✓	r / r	00003.jpg	2009-12-15	2011-03-08	2010-10-24	2901706	0	0	8960518

ASCII ([display](#) - [report](#)) * Hex ([display](#) - [report](#)) * ASCII Strings ([display](#) - [report](#)) * [Export](#) * [View](#) * [Add Note](#)
File Type: JPEG image data, EXIF standard 2.21
Deleted File Recovery Mode

C:/philipshome/Pictures/2010-04-24--11.58.19/00004.jpg

Thumbnail:  [View Full Size Image](#)

NTFS concepts and analysis

- New Technologies File System
 - Microsoft designed and default for MS Windows since Win NT
- Designed for flexibility, scalability, reliability, security and large storage devices
- Key concept of NTFS is that everything in an NTFS system is a file
 - Basic file system administrative data is a file unlike in other systems where it is hidden
 - Entire file system is a data area
 - Only consistent layout is first sectors of the volume contain the boot sector and boot code

NTFS concepts and analysis

- NTFS is a **journaling** file system
 - Changes to the files in the system are entered in a transaction list (a journal) before carrying them out
 - Provides resilience to file system corruption since, in the event of a system failure, the journal entries can be replayed
 - Journal entries stored in attribute \$LogFile
 - \$LogFile can contain forensic artifacts!

Master File Table (MFT)

- Volume Boot Record in sector 0 (of partition)
- The Volume Boot Record (boot sector) has a pointer to the first record of the MFT
- MFT is the heart of NTFS
 - Contains information about all files and directories
- The MFT also has an entry in the MFT

Master File Table (MFT)

- Each file has (at least) one entry in the Master File Table (MFT)
- Each entry in the MFT contains a number of attributes
- Attributes are either resident or non-resident
 - If small enough stored in the MFT (resident)
otherwise a pointer to the location (non-resident)

MFT Entries

- The size of each MFT entry is defined in the boot sector but all versions of Windows use 1024 bytes
- First 42 bytes contain 12 fields
- Remaining 982 bytes used for attributes
- If attribute cannot fit in 1024 bytes then the attribute is marked 'non-resident' and allocated to additional clusters
 - One attribute is the file's contents
 - Usually won't fit in the 1024 bytes

Standard attribute types

- Seventeen standard attributes
- Important ones
 - `$STANDARD_INFORMATION`
 - Flags, dates and times last accessed, owner
 - `$FILE_NAME`
 - File name in unicode
 - `$DATA`
 - File content

Deleting and erasing files

- Regardless of the file system it is surprisingly difficult to completely remove files from a disk
- Most 'delete' operations merely flag the entry in the file table (or equivalent) signifying that the area is available for reuse
- SSD Drives add complexity (Trim)
- Study at MIT in 2003 looked at 158 used disk drives purchased from eBay
 - <http://web.mit.edu/newsoffice/2003/diskdrives.html>
 - Found more than 5,000 credit card numbers, detailed personal and corporate financial records, numerous medical records, gigabytes of personal email and pornography.
 - Only 12 were properly sanitized

Deleting and erasing files

- Deleting and erasing quite different
 - Deleted does not guarantee erasure of the data
- To erase data need to use a utility such as scrub (in Unix systems)
 - Writes a random pattern of 0s and 1s to the disk
 - Does it several times, with newer hard drives 1 wipe is sufficient
 - An interesting paper on erasing data from solid state devices at http://www.usenix.org/events/fast11/tech/full_papers/Wei.pdf
- Zero Only
 - Single pass across drive writing zero values across each sector
- US Airforce standard 5020
 - First fills the drive with binary zeros, then binary ones, and finally an arbitrary character.

Deleting and erasing files

- Australia Signals Directorate (ASD) requires physical destruction for Secret+ data
- ASD standards require 3mm or 6mm pieces depending on sensitivity (e.g. Top Secret)

Hard Drive Destruction



Deleting and Erasing Files for the Paranoid



Image Source: <http://eecue.com/c/driveslag>

Pattern Matching

- Simplest level is matching particular character strings
 - Already seen this in lab 2
 - Made use of piping of text into `grep`
- ```
xxd image.dd | grep secret
```
- ```
strings --encoding=l --radix=d image.dd | grep "\.jpg$"
```
- ```
cat inputfile | grep phil*
```
- Unix has a rich set of pattern matching operations
  - Regular expressions
  - Lots of wild cards and options

# Pattern matching – regular expressions

| Meta character | Description                                                                                                                        |
|----------------|------------------------------------------------------------------------------------------------------------------------------------|
| .              | Matches any single character.<br>Eg pic.hols Will match pic1hols, pic2hols and pic3hols                                            |
| [ ]            | Matches a single character contained within the brackets pic[12]hols will match pic1hols, pic2hols but not pic3hols                |
| [ ^ ]          | Matches a single character that is not contained within the brackets pic[^12]hols will match pic3hols but not pic1hols or pic2hols |
| \$             | Matches the last character of the string<br>Eg pic1hol\$ will match pic1hols                                                       |
| *              | Matches preceding element 0 or more times eg ab*c will match ac, abc, abbbc                                                        |

# Pattern Matching

- Pattern matching not just restricted to alpha-numeric strings
  - Image matching an important area
  - Facial pattern matching
- Not necessarily matching the same image but similar images
  - Usually do not match pixel for pixel but match some attributes of the image
  - Example, facial recognition matches relative position of ears, nose, eyes, chin and mouth
    - Makes use of mathematical transforms such as Fourier, Discrete Cosine , Wavelets and similar

# Summary

- Major file systems
- Detailed look at FAT and NTFS
- A little on pattern matching